


# Intro to Convolutional Neural Networks

BUDSA – Workshop  
By: Christina + Ryan

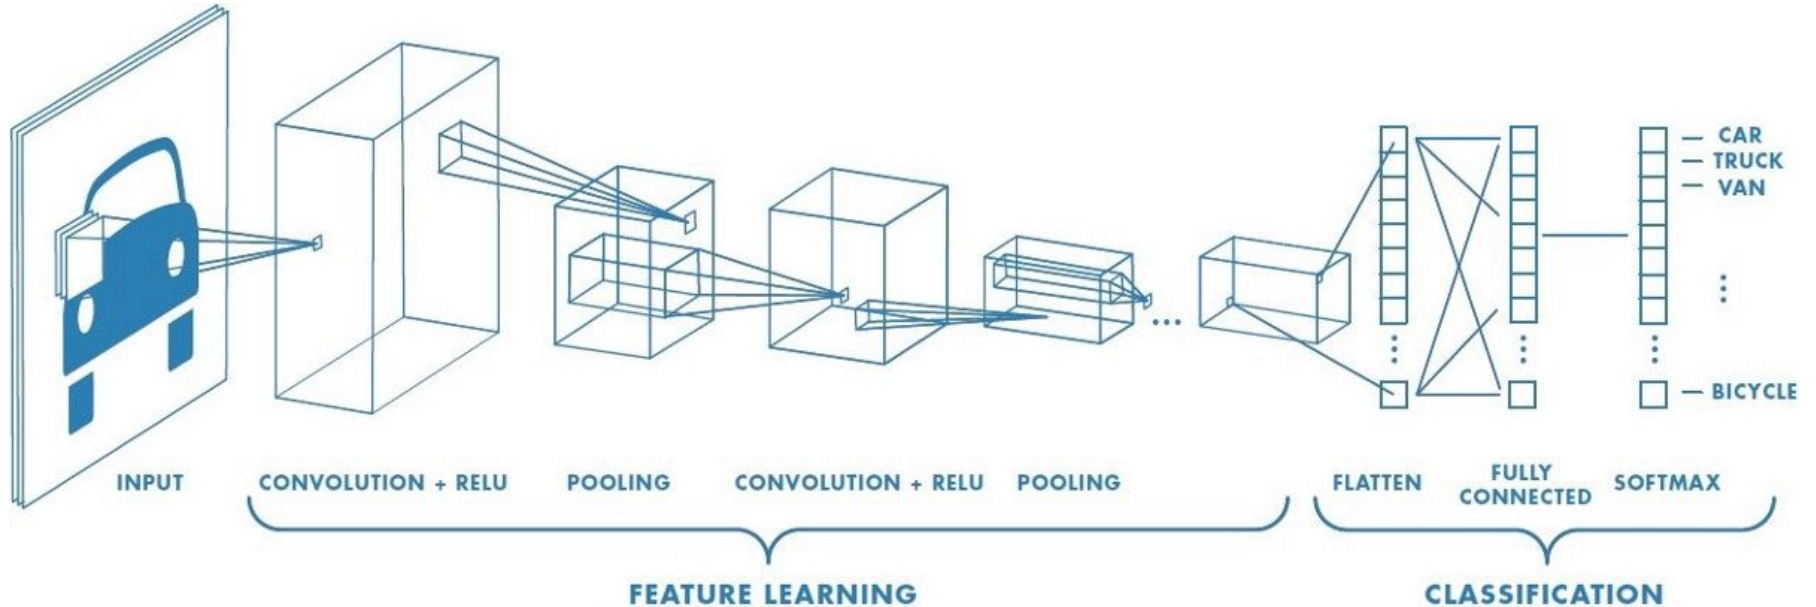


# Sample Architecture

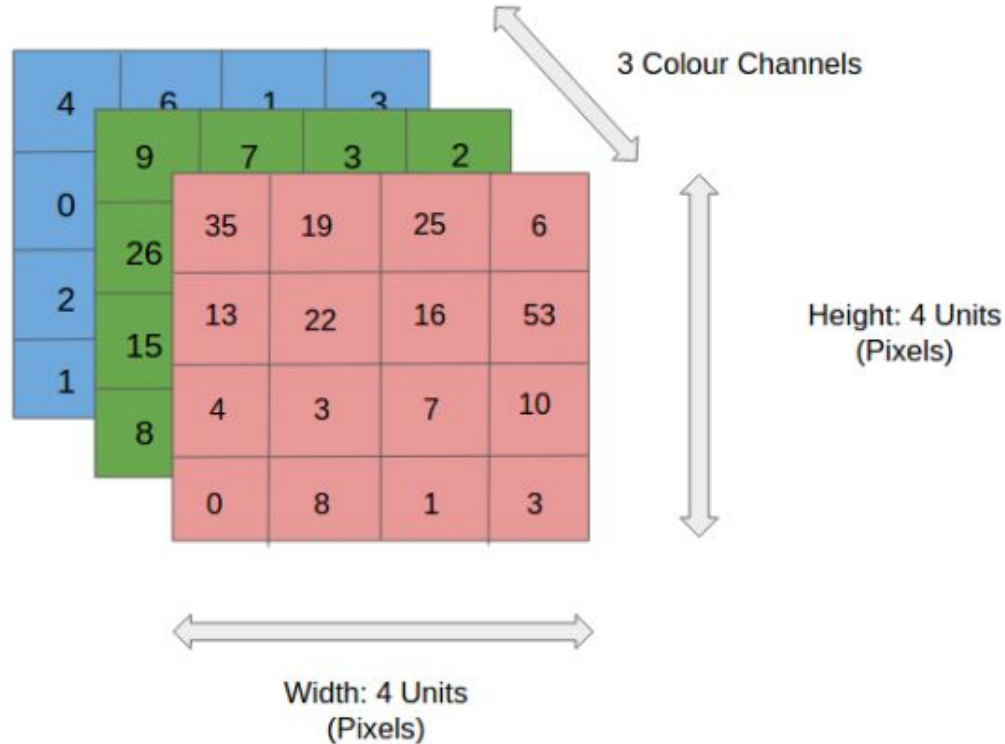
## Convolutional Neural Network




# Sample Architecture



# How do images work?




# Kernels (Filters)



# 1. Convolutional Layer and Nonlinear Transformation

- Preserves locality
- Finds different kernels (filters)
  - Best weights and biases via gradient descent
- Nonlinear transformation functions: ReLU, Tanh, GeLU

Imagine a special projector onto a screen



## 2. Pooling Layer


Ex. Max2D Pooling – accentuates prevalent features within the image

---

# 3. Flatten and Deep Neural Network

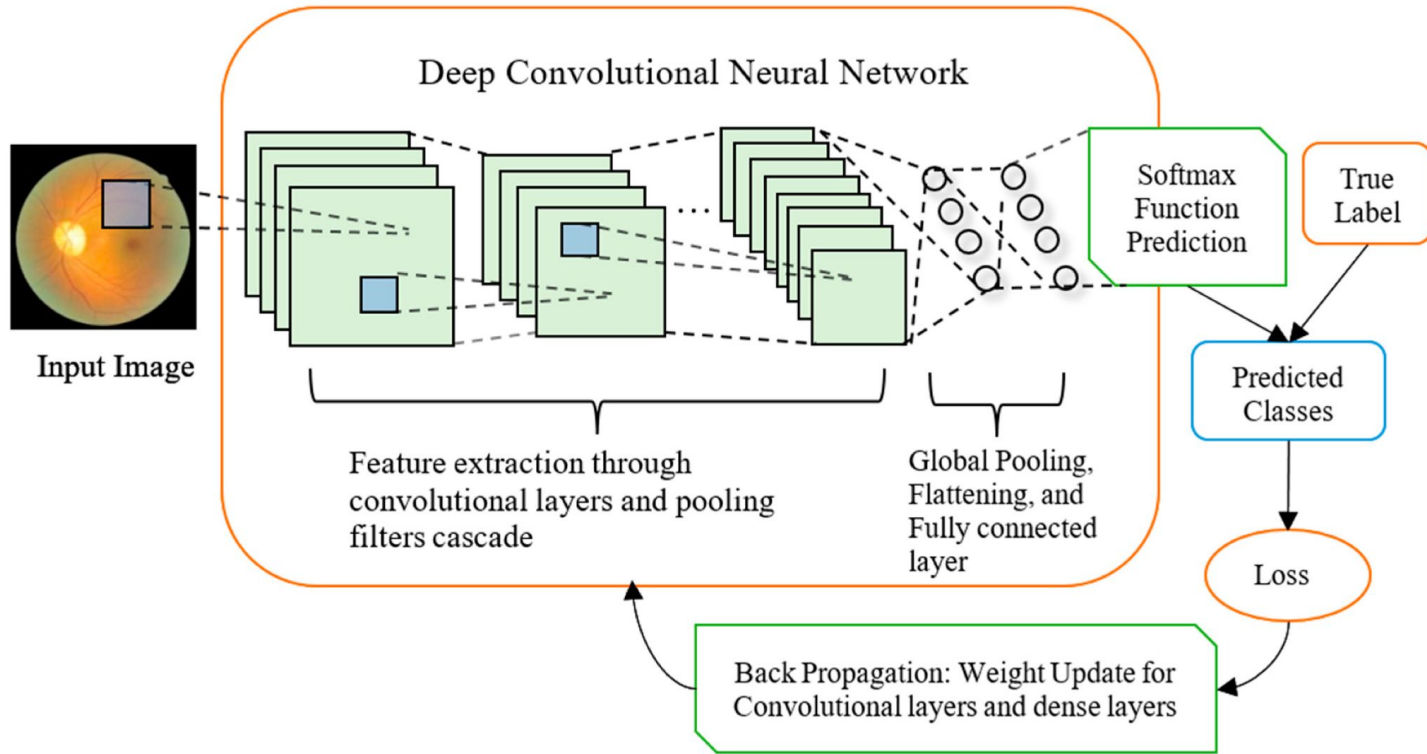
- Flatten the final image into a singular vector
- Fully-connected neural network
- Last layer -- softmax for probabilities of classes, or replace with one neuron for class prediction

## ARTIFICIAL NEURAL NETWORK





# Putting it together: Construction of Architecture




# How to train?

When you create the architecture, you initialize the architecture with some weights and biases

**Backpropagation** is always the answer

- Finds the best kernels (filters)
- Finds the best weights and biases for fully-connected layers



# Why do CNN work so well for image processing?

- Convolution preserves locality
- Max2DPooling layers accentuates features
- Nonlinear transformations allow network to look for latent features (unobservable features of ex. cat/dog)

## Cons?

- Long to train with limited computing power
- Typically need to augment data to get better results
- Vision Transformers are better for SOTA

# Resources

- CNN: <https://poloclub.github.io/cnn-explainer/>
- CNN Cheatsheet: <https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-convolutional-neural-networks>
- Visualizing Kernels: <https://setosa.io/ev/image-kernels/>
- Statquest on CNN: <https://www.youtube.com/watch?v=HGwBXDKFk9I>